

Dear Podman mentorship team,

I first learned about the LFX Mentorship program through a post in Taiwan's open-source community, and this project stood out immediately because it describes work I have already been doing by hand.

I'm interested in this program for a specific reason: I know exactly what the manual version of this job feels like. As a KubeRay contributor, I diagnosed and fixed a flaky e2e test that was intermittently failing CI with 409 Conflict errors — which meant digging through GitHub Actions logs, distinguishing a genuine concurrent-update race condition from infrastructure noise, and shipping a fix that was merged upstream. That triage process was slow, repetitive, and entirely manual. Building an agentic system that does the extraction, categorization, and reporting automatically is the tool I wished I had, and I'd like to help build it for Podman's maintainers.

My experience maps to each part of the proposed toolchain. For the data ingestion side: I've authored a GitHub Actions workflow from scratch and worked with the GitHub API directly, and I'm comfortable in Go, Python, and Shell. For the agentic analysis side: I co-built an LLM pipeline for database schema name expansion, where I designed the evaluation harness — curated gold-label datasets scored with precision/recall and fuzzy matching — because an agent's output is only useful if you can verify it's correct, which is exactly the discipline flake categorization will need. I've also shipped agent-steering documentation before: one of my five merged Apache Airflow PRs was the Traditional Chinese locale guideline for Airflow's AI-assisted translation agent, the same kind of "how maintainers tweak the agent's behavior" document this project lists as an expected outcome.

What I hope to get out of it: I've been contributing to open source self-directed — finding issues, fixing them, learning from review comments. What I haven't had is a maintainer's mentorship on a project with a defined scope, from architecture through deployment and handoff. I want to learn how experienced maintainers think about building tools other maintainers will actually adopt, get properly deep into agentic AI frameworks rather than just consuming LLM APIs, and come out of the three months as a lasting contributor to the Podman ecosystem, not just a term participant.

Thank you for considering my application.

James (Ying-Fang) Jaw